14
months

3
champions

0
purchase orders

The most common outcome for early-stage EdTech selling into higher education.

Cracking higher ed: Why startups miss the mark

Philippos Savvides

Head, ScaleU | ASU's EdTech Accelerator

Not product failure. Procurement failure.

The higher ed paradox

Institutions require proof of enterprise deployment before purchasing

→

Startups cannot demonstrate enterprise capability without institutional access

12-18

month avg. sales cycle

$2.4B

EdTech VC in 2024 (lowest since 2014)

60%

of EdTech startups fail within 5 years

You get one shot. Generate the right evidence, from the right people, fast enough to matter.

The student journey: where products live

Every EdTech product maps to a stage in the learner lifecycle. Your stage determines your buyer, your evidence, and your competition.

Pre-enroll	Apply	Onboard	Select & enroll	Course experience	Graduate & beyond
Readiness, evaluation	Admission, transcripts	Orientation, credit eval	Advising, scheduling	Learning, assessment	Career, alumni

Underserved: high pain, low coverage

Saturated: 15+ products

80% of founders build for Course Experience. The biggest opportunities are everywhere else.

Reframe: market = people + job

People hire products to make progress. The job is what they need to accomplish. Your product is what they hire to do it.

The person	The job	What they hire
Prospective student (Pre-enroll)	Understand if my credits transfer before I commit	**Transfer credit evaluation tool**
Academic advisor (Select & enroll)	Guide 300 students to the right courses each term	**Degree planning intelligence**
Faculty member (Course experience)	Give meaningful feedback to 200 students per section	**AI grading with learning diagnostics**
Career services director (Graduate & beyond)	Connect graduates to employers at scale	**Talent pipeline platform**

If you cannot name the person, the job, and the journey phase, you do not have a market. You have a hypothesis.

Same product, three different jobs

A retention analytics tool. One product. Three stakeholders hiring it for three different reasons.

Academic advisor

Faster triage

"Show me which of my 300 students need intervention this week, ranked by urgency."

VP of Student Affairs

Reduce DFW rates

"Give me a 15% reduction in DFW rates this academic year so I can justify my budget."

Provost

Strategic plan metrics

"I need a retention number for the board presentation that shows year-over-year improvement."

The question: who do you validate with? The advisor's 'yes' does not equal the Provost's budget.

Who do you validate with?

Noise		Signal
"Faculty love it" Faculty don't control budget	→	**Budget holder named a line item** Money is allocated, not just interest
"10,000 student signups" Usage without outcome data	→	**Pilot measured buyer outcomes** Evidence the decision-maker needs
"The dean said transformative" Enthusiasm without procurement	→	**Procurement started before pilot ended** Process is moving, not just praise
"3 pilot champions" None own procurement authority	→	**Product survives if champion leaves** Institutional need, not personal favor

Noise comes from users and champions. Signal comes from buyers and processes.

Three questions before you build

1

What job?

Who is struggling? What is their struggling moment? Is there a budget line item for this pain?

If they haven't tried to solve it, the job isn't urgent enough to pay for.

2

What proof?

What measurable outcome will convince the buyer (not the user) to issue a purchase order?

If the only outcome is 'users liked it,' you don't have evidence.

3

What moat?

Why can't the university's IT team replicate this with a general-purpose LLM and a 2-week sprint?

If AI is your feature, AI is also your substitution risk.

What job, what proof, what moat. If you can't answer all three, you're not ready for a pilot.

The moat problem: AI changes the math

High exposure	Moderate exposure	Low exposure
Content generation, summarization, Q&A, template output	AI embedded in workflow or dataset with switching costs	Structural moat that AI substitution cannot dissolve
Replicable in 12-18 months	**Replicable in 2-4 years**	**Defensible for 3+ years**

Where is your product?

What moves you right?

Proprietary data network

Longitudinal data that improves with usage and cannot be reconstructed by a new entrant

Deep integration

LTI/SIS write-back creating operational switching costs that exceed the product's price

Supply-side network effects

More providers, mentors, or contributors increase quality for all existing users

Regulated access

FERPA, HIPAA, or credentialing compliance that LLM wrappers cannot clear

Lead with the institutional dependency you create, not the technology you use.

The diagnostic: 5 questions before your next pilot

1. **Who is struggling, and what is the struggling moment?**

   If you can't name a specific person with a specific problem, you have a hypothesis, not a job.

2. **What have they already tried (and why did it fail)?**

   Your competition isn't other EdTech. It's spreadsheets, grad students, and ignoring the problem.

3. **Is there a budget line item?**

   "Interesting" vs. "funded" is the only filter that matters.

4. **What happens if they do nothing?**

   "We lose accreditation" is a job. "Students might learn better" is a hypothesis.

5. **Who else must say yes?**

   If the answer is "just me," the pilot succeeds and procurement fails.

Five questions. If you can't answer all five, go back to discovery.

The framework

1. Map your product to a journey phase (and check if it's saturated)
2. Define the job for the person, not the institution
3. Identify the buyer, not just the user
4. Design the pilot to produce evidence the buyer needs
5. Build the moat that AI can't replicate

Philippos Savvides

Head, ScaleU
ASU's EdTech Accelerator

philippos.savvides@asuep.org
scaleu.org
linkedin.com/in/savvides | x.com/savvides

Questions?